

List of References

Connor H. McCurley

Contents

1	Army	2
2	Manifold/ Representation Learning	2
2.1	Classic Methods	2
2.2	Competitive Hebbian Learning	3
3	Manifold Regularization	7
4	Multiple Instance Learning	7
5	Fusion	7
6	Outlier/ Adversarial Detection	7
7	Segmentation	7

1 Army

Hall et al. (2018) - *Probabilistic Object Detection: Definition and Evaluation*

Summary: A probabilistic object detection metric (PDQ - Probability-based Detection Quality) was proposed, thus defining the new task of defining probabilistic object detection metrics. The ability of deep CNNs to quantify both *epistemic* and *aleatoric uncertainty* is paramount for deployment safety-critical applications. PDQ aims to measure the accuracy of an image object detector in terms of its label uncertainty and spatial quality. This is achieved through two steps. First, a detector must reliably quantify its *semantic uncertainty* by providing full probability distributions over known classes for each detection. Next, the detectors must quantify spatial uncertainty by reporting *probabilistic bounding boxes*, where the box corners are modeled as normally distributed. A loss function was constructed to consider both label and spatial quality when providing a final detection measure. The primary benefit of this method is that it provides a measure for the level of uncertainty in a detection.

Is it possible to replace the probabilistic metric with a possibilistic one? Could this be more effective at handling outlying cases?

Mahalanobis and McIntosh (2019) - *A comparison of target detection algorithms using DSIAC ATR algorithm development data set*

Summary: The authors provided an initial characterization of detection performance on the DSIAC dataset using the *Faster R-CNN* algorithm and *Quadratic Correlation Filter (QCF)*. Performance was evaluated on two datasets, “easy” and “difficult”, where the difficulty was determined by number of pixels on target and local contrast. Under difficult conditions, the Faster R-CNN algorithm achieved noteworthy performance, detecting as much as 80% of the targets at a low false alarm rate of 0.01 FA/Square degree. The dataset was limited by a lack of background diversity.

Tanner and Mahalanobis (2019) - *Fundamentals of Target Classification Using Deep Learning*

Summary: A shallow CNN was utilized for ATR on the DSIAC MWIR dataset. The goal of the study was to determine the range of optimal thresholds which would optimally separate the target and clutter class distributions defined by the CNN predictions (output of softmax), as well as determine an upper bound on the number of training images required for optimizing performance. The shallow CNN (5 layers) and a Difference of Gaussians (DoG), which finds regions of high intensity on dark backgrounds were used to detect and classify targets. The CNN could correctly classify 96% of targets as targets and as few as 4% of clutter as targets. It was found that the DoG detector failed when the targets were small (long range) or if the overall image was bright (infrared taken during the daytime). It was also determined that guessing the bright pixels were at the center of the targets was a bad assumption. (The brightest part of a target is not necessarily at its center.)

2 Manifold/ Representation Learning

2.1 Classic Methods

van der Maaten et al. (2007) - *Dimensionality Reduction: A Comparative Review*

Summary:

Jindal and Kumar (2017) - *A Review on Dimensionality Reduction Techniques*

Summary:

Bengio et al. (2012) - *Unsupervised Feature Learning and Deep Learning: A Review and New Perspectives*

Summary:

Tenenbaum et al. (2000) - *A Global Geometric Framework for Nonlinear Dimensionality Reduction*

[Summary:](#)

Roweis and Saul (2000) - *Nonlinear Dimensionality Reduction by Locally Linear Embedding*

[Summary:](#)

Saul and Roweis (2001) - *An introduction to locally linear embedding*

[Summary:](#)

Belkin and Niyogi (2003) - *Laplacian Eigenmaps for Dimensionality Reduction and Data Representation*

[Summary:](#)

Bishop et al. (1998) - *GTM: The Generative Topographic Mapping*

[Summary:](#)

Delaporte et al. (2008) - *An introduction to diffusion maps*

[Summary:](#)

Theodoridis and Koutroumbas (2008b) - *The Karhunen-Loeve Transform*

[Summary:](#)

Theodoridis and Koutroumbas (2008a) - *Kernel PCA*

[Summary:](#)

Tipping and Bishop (1999) - *Probabilistic Principal Component Analysis*

[Summary:](#)

Lawrence (2003) - *Gaussian Process Latent Variable Models for Visualisation of High Dimensional Data*

[Summary:](#)

Lawrence (2005) - *Probabilistic Non-linear Principal Component Analysis with Gaussian Process Latent Variable Models*

[Summary:](#)

Gorban and Zinovyev (2008) - *Elastic Maps and Nets for Approximating Principal Manifolds and Their Application to Microarray Data Visualization*

[Summary:](#)

Lee et al. (2016) - *Learning Representations from Multiple Manifolds*

[Summary:](#)

2.2 Competitive Hebbian Learning

Rumelhart and Zipser (1985) - *Feature Discovery by Competitive Learning*

[Summary:](#)

Kohonen (1990) - *The self-organizing map*

Summary: The self-organizing map (SOM) creates spatially organized intrinsic representations of features. It belongs to the category of neural networks which use “competitive learning”, or “self-organization”. It is a sheet-like artificial neural network in which the cells become tuned to various input patterns through an unsupervised learning process. Only a neighborhood of cells give an active response to the current input sample. The spatial location or coordinates of cells in the network correspond to different modes of the input distribution. The self-organizing map is also a form of vector quantization (VQ). The purpose of VQ is to approximate a continuous probability density function $p(\mathbf{x})$ of input vectors \mathbf{x} using a finite number of codebook vectors, \mathbf{m}_i , $i = 1, 2, \dots, k$. After the “codebook” is chosen, the approximation of \mathbf{x} involves finding the reference vector, \mathbf{m}_c closest to \mathbf{x} . The “winning” codebook vector for sample \mathbf{x} satisfies the following:

$$\|\mathbf{x} - \mathbf{m}_c\| = \min_i \|\mathbf{x} - \mathbf{m}_i\|$$

The algorithm operates by first initializing a spatial lattice of codebook elements (also called “units”), where each unit’s representative is in $\mathbf{m}_i \in \mathbb{R}^D$ where D is the dimensionality of the input samples \mathbf{x} . The training process proceeds as follows. A random sample is selected and presented to the network and each unit determines its activation by computing dissimilarity. The unit who’s codebook vector provides the smallest dissimilarity is referred to as the *winner*.

$$c(t) = \arg \min_i d(\mathbf{x}(t), \mathbf{m}_i(t))$$

Both the winning vector and all vectors within a neighborhood of the winner are updated toward the sample by

$$\mathbf{m}_i(t+1) = \mathbf{m}_i(t) + \alpha(t) \cdot h_{ci}(t) \cdot [\mathbf{x}(t) - \mathbf{m}_i(t)]$$

where $\alpha(t)$ is a learning rate which decreases over time and $h_{ci}(t)$ is a neighborhood function which is typically unimodal and symmetric around the location of the winner which monotonically decreases with increasing distance from the winner. A radial basis kernel is typically chosen for the neighborhood function as

$$h_{ci}(t) = \exp\left(-\frac{\|\mathbf{r}_c - \mathbf{r}_i\|^2}{2\sigma^2(t)}\right)$$

where the top expression represents the Euclidean distance between units c and i with \mathbf{r}_i representing the 2-D location of unit i in the lattice. The neighborhood kernel’s bandwidth is typically initialized to a value which covers a majority of the input space and decreases over time such that solely the winner is adapted toward the end of the training procedure.

The SOM essentially performs density estimation of high-dimensional data and represents it in a 2 or 3-D representation. At test time, the dissimilarity between each unit in the map and an input sample are computed. This dissimilarity can be used to effectively detect outliers, thus making the SOM a robust method which can provide confidence values for it’s representation abilities.

In this paper, the SOM was applied to speech recognition, but made note of previous uses in robotics, control of diffusion processes, optimization problems, adaptive telecommunications, image compression, sentence understanding, and radar classification of sea-ice.

Rauber et al. (2002) - *The growing hierarchical self-organizing map: exploratory analysis of high-dimensional data*

Summary: The Growing Hierarchical Self-organizing Map (GHSOM) is an extension of the classical SOM. It is an artificial neural network with a hierarchical architecture, composed of individually growing SOMs. Layer 0 is composed of a single neuron representing the mean of the training data. A global stopping criteria is developed as a fraction of the mean quantization error. This means that all units must represent their respective subsets of data an a MQE smaller than a fraction of the 0 layer mean quantization error. For all units not satisfying this criteria, more representation is required for that area of the feature space and additional units are added. After a particular number of training iterations, the quantization errors are computed and the unit with the highest error is selected as the *error unit*. The most dissimilar neighbor of the error unit is chosen and a row/ column of nodes is injected between them. The growth process continues until a second stopping criteria is met. Any units still not satisfying the global criteria are deemed to need extra representation. Child map are initialized below these units and trained with the subset of data mapped to its parent node.

In conclusion, the GHSOM is a growing self-organizing map architecture which has the ability to grow itself until the feature space is adequately represented. For areas of the space needing a more specific level of granularity, a hierarchical structure is imposed to “fill-in” areas of high density.

The GHSOM has been applied to the areas of finance, computer network traffic analysis, manufacturing and image analysis (Palomo 2017).

Chiang and Gader (1997) - *Hybrid fuzzy-neural systems in handwritten word recognition*

Summary:

Frigui and Gader (2009) - *Detection and Discrimination of Land Mines in Ground-Penetrating Radar Based on Edge Histogram Descriptors and a Possibilistic K-Nearest Neighbor Classifier*

Summary:

Fritzke (1994) - *A Growing Neural Gas Network Learns Topologies*

Summary: Abstract: An incremental network model is introduced which is able to learn the important topological relations in a given set of input vectors by means of a simple Hebb-like learning rule. In contrast to previous approaches like the “neural gas” method of Martinetz and Schulten (1991, 1994), this model has no parameters which change over time and is able to continue learning, adding units and connections, until a performance criterion has been met. Applications of the model include vector quantization, clustering, and interpolation.

In contrast to SOMs and “growing cell structures”, which can project data onto non-linear subspaces which are chosen *a priori*, the GNG is able to adapt its topology to match that of the input data distribution. The growing process continues until a pre-defined level of quantization error has been reached.

The base algorithm is outlined in Palomo (2017), *Growing Hierarchical Neural Gas Self-Organizing Network*.

Palomo and Lopez-Rubio (2017) - *The Growing Hierarchical Neural Gas Self-Organizing Neural Network*

Summary: Abstract: The growing neural gas (GNG) self-organizing neural network stands as one of the most successful examples of unsupervised learning of a graph of processing units. Despite its success, little attention has been devoted to its extension to a hierarchical model, unlike other models such as the self-organizing map, which has many hierarchical versions. Here, a hierarchical GNG is presented, which is designed to learn a tree of graphs. Moreover, the original GNG algorithm is improved by a distinction between a growth phase where more units are added until no significant improvement in the quantization error is obtained, and a convergence phase where no unit creation is allowed. This means that a principled mechanism is established to control the growth of the structure. Experiments are reported, which demonstrate the self-organization and hierarchy learning abilities of our approach and its performance for vector quantization applications. Experiments were performed in structure learning, color quantization, and video sequence clustering.

The aim of this method was to improve the adaptation ability of the Growing Hierarchical Self-Organizing Map proposed by Rauber (2002). This was to be done through the extension of the Growing Neural Gas, which disposes of the fixed lattice topology enforced by the SOM. Additionally, the GNG learns a dynamic graph with variable numbers of neurons and connections. The graph represents the input data in a more plastic and flexible way than the fixed-topology map.

All clustering methods that learn a hierarchical structure have advantages even when used for non-hierarchical data. The learned hierarchical structure can be pruned at several levels, which yields alternative representations of the input data set at different levels of detail. This can be used to visualize a data set in coarser or more detailed way. For vector quantization applications, the different pruning levels correspond to smaller or larger codebooks, so that a balance can be attained between the size of the codebook and the quantization error within the same hierarchical structure.

The growing hierarchical neural gas (GHNG) model is defined as a tree of self-organizing graphs. Each graph is made of a variable number of neurons or processing units, so that its size can grow or shrink during learning. In addition, each graph is the child of a unit in the upper level, except for the top level (root) graph. The training procedure is described by the following:

Each graph begins with $H \geq 2$ units and one or more undirected connections between them. Both the units and connections can be created and destroyed during the learning process. It is also not necessary that the graph is connected. Let the training set be denoted as \mathcal{S} with $\mathcal{S} \subset \mathbb{R}^D$, where D is the dimensionality of the input space. Each unit $i \in \{1, \dots, H\}$ has an associated prototype $\mathbf{w}_i \in \mathbb{R}^D$ and an error variable $e_i \in \mathbb{R}$, $e_i \geq 0$. Each connection has an associated age, which is a nonnegative integer. The set of connections will be notetd as $A \subseteq \{1, \dots, H\} \times \{1, \dots, H\}$. The learning mechanism for the GHNG is based on the original GNG, but includes a novel procedure to control the growth of the graph. First, a growth phase is performed where the graph is allowed to enlarge until a condition is met, which indicates that further growing would provide no significant improvement in the quantization error. After that, a convergence phase is executed where no unit creation is allowed in order to carry out a fine tuning of the graph. the leraning algorithm is provided in the following steps.

I believe the author’s experimental approach did not take advantage of the method’s strengths. The author’s only demonstrated experiments in vector quantization, and used corresponding metrics. This method could be used to represent manifold topology of differing dimensionality. This could be useful in HSI imagery, for example where different environment patches require manifold representations of various dimensionality. Additionally, this could potentially be used to handle the sensor fusion problem with sensor loss/ drop-out.

Sun et al. (2017) - *Online growing neural gas for anomaly detection in changing surveillance scenes*
Summary:

Lopez-Rubio and Palomo (2011) - *Growing Hierarchical Probabilistic Self-Organizing Graphs*
Summary:

Palomo and Lopez-Rubio (2016) - *Learning Topologies with the Growing Neural Forest*
Summary:

3 Manifold Regularization

Tsang and Kwok (2007) - *Large-Scale Sparsified Manifold Regularization*
Summary:

Ren et al. (2017) - *Unsupervised Classification of Polarimetric SAR Image Via Improved Manifold Regularized Low-Rank Representation With Multiple Features*
Summary:

Belkin et al. (2006) - *Manifold Regularization: A Geometric Framework for Learning from Labeled and Unlabeled Examples*
Summary:

Ratlle et al. (2010) - *Semisupervised Neural Networks for Efficient Hyperspectral Image Classification*
Summary:

Li et al. (2015) - *Approximate Policy Iteration with Unsupervised Feature Learning based on Manifold Regularization*
Summary:

Meng and Zhan (2018) - *Zero-Shot Learning via Low-Rank-Representation Based Manifold Regularization*
Summary:

4 Multiple Instance Learning

-

Summary:

-

Summary:

5 Fusion

6 Outlier/ Adversarial Detection

7 Segmentation

Caselles et al. (1997) - *Geodesic Active Contours*
Summary:

Álvarez et al. (2010) - *Morphological Snakes*

Summary: The authors introduce a morphological approach to curve evolution. Snakes or curves iteratively solve partial differential equations (PDEs). By doing so, the shape of the snake deforms to minimize the internal and external energies along its boundary. The internal component keeps the curve smooth, while the external component attaches the curve to image structures such as edges, lines, etc. Curve evolution is one of the most widely used image segmentation/ object tracking algorithms. The main contribution of the paper is a new morphological approach to the solution of the PDE associated with snake model evolution. They approach the solution using only inf-sup operators which has the main benefit of providing simpler level sets (0 outside the contours and 1 inside).

Márquez-Neila et al. (2014) - *A Morphological Approach to Curvature-Based Evolution of Curves and Surfaces*

Summary:

References

- L. Álvarez, L. Baumela, P. Henríquez, and P. Márquez-Neila. Morphological snakes. In *2010 IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, pages 2197–2202, June 2010. doi: 10.1109/CVPR.2010.5539900.
- M. Belkin and P. Niyogi. Laplacian eigenmaps for dimensionality reduction and data representation. *Neural Computation*, 15(6):1373–1396, 2003. doi: 10.1162/089976603321780317.
- M. Belkin, P. Niyogi, and V. Sindhwani. Manifold regularization: A geometric framework for learning from labeled and unlabeled examples. *J. Mach. Learn. Res.*, 7:2399–2434, December 2006. ISSN 1532-4435. URL <http://dl.acm.org/citation.cfm?id=1248547.1248632>.
- Y. Bengio, A. C. Courville, and P. Vincent. Unsupervised feature learning and deep learning: A review and new perspectives. *CoRR*, abs/1206.5538, 2012. URL <http://arxiv.org/abs/1206.5538>.
- C. Bishop, M. Svensn, and C. K. I. Williams. Gtm: The generative topographic mapping. 10:215–234, January 1998. URL <https://www.microsoft.com/en-us/research/publication/gtm-the-generative-topographic-mapping/>.
- V. Caselles, R. Kimmel, and G. Sapiro. Geodesic active contours. *International Journal of Computer Vision*, 22(1):61–79, Feb 1997. ISSN 1573-1405. doi: 10.1023/A:1007979827043. URL <https://doi.org/10.1023/A:1007979827043>.
- J. Chiang and P. D. Gader. Hybrid fuzzy-neural systems in handwritten word recognition. *IEEE Transactions on Fuzzy Systems*, 5(4):497–510, Nov 1997. ISSN 1063-6706. doi: 10.1109/91.649901.
- J. Delaporte, B. M. Herbst, W. Hereman, and S. Van der Walt. An introduction to diffusion maps. 2008.
- H. Frigui and P. Gader. Detection and discrimination of land mines in ground-penetrating radar based on edge histogram descriptors and a possibilistic k -nearest neighbor classifier. *IEEE Transactions on Fuzzy Systems*, 17(1), Feb 2009. ISSN 1063-6706. doi: 10.1109/TFUZZ.2008.2005249.
- B. Fritzke. A growing neural gas network learns topologies. In *Proceedings of the 7th International Conference on Neural Information Processing Systems, NIPS’94*, pages 625–632, Cambridge, MA, USA, 1994. MIT Press. URL <http://dl.acm.org/citation.cfm?id=2998687.2998765>.
- A. N. Gorban and A. Y. Zinovyev. Elastic maps and nets for approximating principal manifolds and their application to microarray data visualization. In A. N. Gorban, B. Kégl, D. C. Wunsch, and A. Y. Zinovyev, editors, *Principal Manifolds for Data Visualization and Dimension Reduction*, pages 96–130, Berlin, Heidelberg, 2008. Springer Berlin Heidelberg. ISBN 978-3-540-73750-6.
- D. Hall, F. Dayoub, J. Skinner, P. Corke, G. Carneiro, and N. Sünderhauf. Probability-based detection quality (PDQ): A probabilistic approach to detection evaluation. *CoRR*, abs/1811.10800, 2018. URL <http://arxiv.org/abs/1811.10800>.
- P. Jindal and D. Kumar. A review on dimensionality reduction techniques. *International Journal of Computer Applications*, 173, 09 2017.
- T. Kohonen. The self-organizing map. *Proceedings of the IEEE*, 78(9):1464–1480, Sep. 1990. ISSN 0018-9219. doi: 10.1109/5.58325.
- N. Lawrence. Probabilistic non-linear principal component analysis with gaussian process latent variable models. *J. Mach. Learn. Res.*, 6:1783–1816, dec 2005. ISSN 1532-4435. URL <http://dl.acm.org/citation.cfm?id=1046920.1194904>.
- N. D. Lawrence. Gaussian process latent variable models for visualisation of high dimensional data. In *Proceedings of the 16th International Conference on Neural Information Processing Systems, NIPS’03*, pages 329–336, Cambridge, MA, USA, 2003. MIT Press. URL <http://dl.acm.org/citation.cfm?id=2981345.2981387>.

- C. Lee, A. Elgammal, and M. Torki. Learning representations from multiple manifolds. *Pattern Recogn.*, 50(C):74–87, February 2016. ISSN 0031-3203. doi: 10.1016/j.patcog.2015.08.024. URL <http://dx.doi.org/10.1016/j.patcog.2015.08.024>.
- H. Li, D. Liu, and D. Wang. Approximate policy iteration with unsupervised feature learning based on manifold regularization. In *2015 International Joint Conference on Neural Networks (IJCNN)*, pages 1–6, July 2015.
- E. Lopez-Rubio and E. J. Palomo. Growing hierarchical probabilistic self-organizing graphs. *IEEE Transactions on Neural Networks*, 22(7):997–1008, July 2011. ISSN 1045-9227. doi: 10.1109/TNN.2011.2138159.
- A. Mahalanobis and B. McIntosh. A comparison of target detection algorithms using dsia atr algorithm development data set. *Proc.SPIE*, Apr. 2019.
- P. Márquez-Neila, L. Baumela, and L. Álvarez. A morphological approach to curvature-based evolution of curves and surfaces. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 36(1):2–17, Jan 2014. ISSN 0162-8828.
- M. Meng and X. Zhan. Zero-shot learning via low-rank-representation based manifold regularization. *IEEE Signal Processing Letters*, 25(9):1379–1383, Sep. 2018. ISSN 1070-9908. doi: 10.1109/LSP.2018.2857201.
- E. J. Palomo and E. Lopez-Rubio. Learning topologies with the growing neural forest. *International Journal of Neural Systems*, 26(04):1650019, 2016. doi: 10.1142/S0129065716500192. URL <https://doi.org/10.1142/S0129065716500192>. PMID: 27121995.
- E. J. Palomo and E. Lopez-Rubio. The growing hierarchical neural gas self-organizing neural network. *IEEE Transactions on Neural Networks and Learning Systems*, 28(9):2000–2009, Sep. 2017. ISSN 2162-237X. doi: 10.1109/TNNLS.2016.2570124.
- F. Ratle, G. Camps-Valls, and J. Weston. Semisupervised neural networks for efficient hyperspectral image classification. *IEEE Transactions on Geoscience and Remote Sensing*, 48(5):2271–2282, May 2010. ISSN 0196-2892. doi: 10.1109/TGRS.2009.2037898.
- A. Rauber, D. Merkl, and M. Dittenbach. The growing hierarchical self-organizing map: exploratory analysis of high-dimensional data. *IEEE Transactions on Neural Networks*, 13(6):1331–1341, Nov 2002. ISSN 1045-9227. doi: 10.1109/TNN.2002.804221.
- B. Ren, B. Hou, J. Zhao, and L. Jiao. Unsupervised classification of polarimetric sar image via improved manifold regularized low-rank representation with multiple features. *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing*, 10(2):580–595, Feb 2017. ISSN 1939-1404. doi: 10.1109/JSTARS.2016.2573380.
- S. T. Roweis and L. K. Saul. Nonlinear dimensionality reduction by locally linear embedding. *Science*, 290(5500):2323–2326, 2000. ISSN 0036-8075. doi: 10.1126/science.290.5500.2323. URL <https://science.sciencemag.org/content/290/5500/2323>.
- D. E. Rumelhart and D. Zipser. Feature discovery by competitive learning. *Cognitive Science*, 9(1):75–112, 1985. doi: 10.1207/s15516709cog0901_5. URL https://onlinelibrary.wiley.com/doi/abs/10.1207/s15516709cog0901_5.
- L. K. Saul and S. T. Roweis. An introduction to locally linear embedding. *Journal of Machine Learning Research*, 7, 01 2001.
- Q. Sun, H. Liu, and T. Harada. Online growing neural gas for anomaly detection in changing surveillance scenes. *Pattern Recognition*, 64:187 – 201, 2017. ISSN 0031-3203. doi: <https://doi.org/10.1016/j.patcog.2016.09.016>. URL <http://www.sciencedirect.com/science/article/pii/S0031320316302771>.
- I. L. Tanner and A. Mahalanobis. Fundamentals of target classification using deep learning. *Proc.SPIE*, Apr. 2019.

- J. B. Tenenbaum, V. Silva, and J. C. Langford. A global geometric framework for nonlinear dimensionality reduction. *Science*, 290(5500):2319–2323, 2000. ISSN 0036-8075. doi: 10.1126/science.290.5500.2319. URL <https://science.sciencemag.org/content/290/5500/2319>.
- S. Theodoridis and K. Koutroumbas. Kernel pca. In *Pattern Recognition, Fourth Edition*, chapter 6, pages 351–353. Academic Press, Inc., Orlando, FL, USA, 4th edition, 2008a. ISBN 1597492728, 9781597492720.
- S. Theodoridis and K. Koutroumbas. The karhunen-loeve transform. In *Pattern Recognition, Fourth Edition*, chapter 6, pages 326–334. Academic Press, Inc., Orlando, FL, USA, 4th edition, 2008b. ISBN 1597492728, 9781597492720.
- M. E. Tipping and C. Bishop. Probabilistic principal component analysis. *Journal of the Royal Statistical Society, Series B*, 21(3):611–622, January 1999. URL <https://www.microsoft.com/en-us/research/publication/probabilistic-principal-component-analysis/>.
- I. W. Tsang and J. T. Kwok. Large-scale sparsified manifold regularization. In B. Schölkopf, J. C. Platt, and T. Hoffman, editors, *Advances in Neural Information Processing Systems 19*, pages 1401–1408. MIT Press, 2007. URL <http://papers.nips.cc/paper/3005-large-scale-sparsified-manifold-regularization.pdf>.
- L. van der Maaten, E. Postma, and H. Herik. Dimensionality reduction: A comparative review. *Journal of Machine Learning Research - JMLR*, 10, 01 2007.